

---

**Started on** Saturday, 22 June 2024, 11:47 PM

---

**State** Finished

---

**Completed on** Saturday, 22 June 2024, 11:54 PM

---

**Time taken** 6 mins 22 secs

---

**Question 1**

Correct

Marked out of 2.00

**Task:** Considering **Recurrent Neural Networks** (RNN), select **architectures** that help to increase **long-term dependencies** and mitigate **vanishing gradients**.

Select one or more:

- ☐ a. Rectified Linear Units
- ☒ b. Long Short Term Memory cells ✓
- ☐ c. Inception Modules
- ☐ d. Residual Blocks
- ☒ e. Gated Recurrent Unit cells ✓

Die Antwort ist richtig.

The correct answers are: Long Short Term Memory cells, Gated Recurrent Unit cells

**Question 2**

Partially correct

Marked out of 1.00

Below are statements on word embeddings.

**Task:** Select all correct statements.

- ☒ a. Word embeddings can model the relationship between words. ✓
- ☒ b. Word embedding vectors are typically smaller compared to one-hot encoded word representations. ✓
- ☐ c. When adding one word to the vocabulary, the dimension of the word embedding space grows by one.
- ☐ d. Typically word embedding vectors are very sparse and can be efficiently implemented using indexed representations.
- ☐ e. Words with similar meanings have similar word embedding vectors.
- ☐ f. One-hot encoded word representations are the same as word embeddings.

Your answer is partially correct.

You have correctly selected 2.

The correct answers are:

Word embeddings can model the relationship between words.,

Words with similar meanings have similar word embedding vectors.,

Word embedding vectors are typically smaller compared to one-hot encoded word representations.

**Question 3**

Incorrect

Marked out of 1.00

Consider you want to train a **word embedding** for a **vocabulary of 1000 words**. In order to capture the meanings and relations of all words, your **embedding size** should also amount to **1000**.

**Question:** Is this correct or not?

Select one:

- ☒ True ✗
- ☐ False

The correct answer is 'False'.

**Question 4**

Incorrect

Marked out of 1.00

Consider you want to train a recurrent neural network (RNN) to translate from english to german sentences.

**Question:** Which RNN architecture is most suited for this task?

- ☐ a. Many-to-many (unsynchronized)
- ☒ b. Many-to-many (synchronized) ✖
- ☐ c. Many-to-one
- ☐ d. One-to-many

Your answer is incorrect.

The correct answer is:

Many-to-many (unsynchronized)

**Question 5**

Incorrect

Marked out of 1.00

**Question:** Using **recurrent neural networks** in **encoder-decoder architecture**, what should be used as **initial hidden input** for the **decoder**?

- ☒ a. The last encoder output. ✖
- ☐ b. The average encoder output.
- ☐ c. The last encoder hidden state.
- ☐ d. The last decoder hidden state.
- ☐ e. The sum of all encoder hidden states.
- ☐ f. The first encoder hidden state.
- ☐ g. The average of all encoder hidden states.

Your answer is incorrect.

The correct answer is:

The last encoder hidden state.

**Question 6**

Incorrect

Marked out of 2.00

Consider you successfully trained a **word embedding** on a huge corpus of text.

**Question:** Which of the following equations do you expect to be valid? Select all that apply.

Select one or more:

- ☒ a.  $e_{\text{lounging}} - e_{\text{learning}} \approx e_{\text{pass}} - e_{\text{fail}}$  ✗
- ☐ b.  $e_{\text{uncle}} - e_{\text{man}} \approx e_{\text{aunt}} - e_{\text{woman}}$
- ☐ c.  $e_{\text{fail}} - e_{\text{lounging}} \approx e_{\text{pass}} - e_{\text{learning}}$
- ☐ d.  $e_{\text{uncle}} - e_{\text{man}} \approx e_{\text{woman}} - e_{\text{aunt}}$
- ☒ e.  $e_{\text{uncle}} - e_{\text{aunt}} \approx e_{\text{woman}} - e_{\text{man}}$  ✗

Die Antwort ist falsch.

The correct answers are:  $e_{\text{fail}} - e_{\text{lounging}} \approx e_{\text{pass}} - e_{\text{learning}}$   
,  $e_{\text{uncle}} - e_{\text{man}} \approx e_{\text{aunt}} - e_{\text{woman}}$